

Machine Learning

Assignment 2:-

(1)

Vocab size $\Rightarrow 292$ after preprocessing.

Assume each doc multinomial distribution.

$p(x/y) \rightarrow$ multinomial

estimate $p_c \Rightarrow$ parameter for multinomial distribution

& obtain w_c and β_c from $p(y)$ and $\log p(x/y)$

\rightarrow

p_c is a vector, ~~for~~ ^{each} document for ~~each~~ ^{each} class, with each entry as occurrence of word "i" in all documents of class "j" and normalized with total words in class "j"

$$\therefore p_c = \left[\frac{\text{word 1 count}}{\text{total words in class 'j'}}, \dots, \frac{\text{word n count}}{\text{total words in class 'j'}} \right]$$

$$\therefore p_c = \frac{1}{\sum_{i \in \text{class } j} x_i} \sum_{i \in \text{class } j} x_i \quad \text{given, } x = \text{frequency count, } x = [x_1, x_2, \dots, x_m]$$

Probability of new document "x" in class c is proportional to,

$$p(x) \propto \pi_c \prod_{j=1}^m p_{c,j}^{x_j}$$

preventing ~~under~~ ^{under}flow by taking log,

$$\propto \log \left(\pi_c \prod_{j=1}^m p_{c,j}^{x_j} \right)$$

$$\propto \log \pi_c + \sum_{j=1}^m x_j \log p_{c,j}$$

∴ choosing max of all, we get classifier,

$$\hat{y} = \underset{c}{\operatorname{argmax}} \left(\log \pi_c + \sum_{j=1}^m x_j \log p_{cj} \right)$$

this is of form, $\hat{y} = \underset{c}{\operatorname{argmax}} (w_c^T x + \beta_c)$

$$\boxed{\begin{aligned} w_c(j) &= \log p_{cj} & \beta_c &= \log \pi_c \end{aligned}}$$

(2) $p(x|y) \Rightarrow$ multivariate normal distribution,

covariance matrix same $\Rightarrow \Sigma_1 = \Sigma_2 = \dots = \Sigma_m = \Sigma$

∴ classifier similar to Linear Discriminant Analysis (LDA)
parameters $\Rightarrow \mu_c$ and Σ

→

Since bag-of-words doesn't always follow a normal distribution,
we convert our dataset as a ~~tfidf~~ TF-IDF matrix,

$$\text{TF (term frequency)} = \frac{\# \text{ Term } t \text{ in document } d}{\text{Total terms in document } d}$$

(Importance to more frequent words).

$$\text{IDF (Inverse document Freq)} = \log \left(\frac{\text{Total documents in corpus } D}{\# \text{ documents with term } t} \right)$$

(Reduce impact of high common words like "the", and
adds importance to rare words).

The same is implemented in code,

Linear Discriminant Analysis Model (LDA)

parameters $\Rightarrow \mu_c$ and Σ

$$\mu_c = \frac{1}{n_c} \sum_{i \in \text{class } c} x_i$$

$$\begin{aligned} \Sigma &= \frac{1}{n} \sum_{i=1}^n (x_i - \mu_c) (x_i - \mu_c)^T \\ &= \frac{1}{n} \sum x_i x_i^T - x_i \mu_c^T - \mu_c x_i^T + \mu_c \mu_c^T \\ &= \frac{1}{n} \left[\sum_{i=1}^n x_i x_i^T - \underbrace{\left(\sum_{i=1}^n x_i \right) \mu_c^T}_{\substack{n_c \mu_c \\ \text{common}}} - \underbrace{\sum_{i=1}^n \mu_c x_i^T}_{\substack{\text{common} \\ \mu_c^T \sum_{i=1}^n x_i}} + \sum_{i=1}^n \mu_c \mu_c^T \right] \\ &= \frac{1}{n} \left[\sum_{i=1}^n x_i x_i^T - \cancel{\sum_{i=1}^n x_i \mu_c^T} - \cancel{\sum_{i=1}^n \mu_c x_i^T} + \sum_{i=1}^n \mu_c \mu_c^T \right] \end{aligned}$$

$$\Sigma = \frac{1}{n} \sum_{i=1}^n x_i x_i^T - \frac{1}{n} \sum_{c=1}^k n_c \mu_c \mu_c^T$$

The probability that a doc belongs to class,

$$\pi_c \det(2\pi\Sigma)^{-1/2} \exp\left(-\frac{1}{2} (x - \mu_c)^T \Sigma^{-1} (x - \mu_c)\right)$$

Taking log,

$$\log \left(\underbrace{\pi_c \det(2\pi\Sigma)^{-1/2}}_{\substack{\text{Constant} \\ \text{ignore}}} \exp\left(-\frac{1}{2} (x - \mu_c)^T \Sigma^{-1} (x - \mu_c)\right) \right)$$

$$\therefore = \log \pi_c + \left(-\frac{1}{2} (x - \mu_c)^T \Sigma^{-1} (x - \mu_c) \right)$$

$$= \log \pi_c + \left(-\frac{1}{2} \left(\underbrace{x^T \Sigma^{-1} x}_{\substack{\text{Constant} \\ \text{term}}} - x^T \Sigma^{-1} \mu_c - \mu_c^T \Sigma^{-1} x + \mu_c^T \Sigma^{-1} \mu_c \right) \right)$$

Choosing maximum of all classes gives us the classifier,

$$\hat{y} = \underset{c}{\operatorname{argmax}} \left(\log \pi_c - \frac{1}{2} \mu_c^T \Sigma^{-1} \mu_c + \mu_c^T \Sigma^{-1} x \right)$$

$$\left[\sum_{i=1}^n \mu_c^T \Sigma^{-1} \mu_c - \sum_{i=1}^n \mu_c^T \Sigma^{-1} x_i \right] \frac{1}{n}$$

$$\left(\sum_{i=1}^n \mu_c^T \Sigma^{-1} \mu_c - \sum_{i=1}^n \mu_c^T \Sigma^{-1} x_i \right) \frac{1}{n}$$

$$\left[\sum_{i=1}^n \mu_c^T \Sigma^{-1} \mu_c - \sum_{i=1}^n \mu_c^T \Sigma^{-1} x_i \right] \frac{1}{n}$$

$$\left((U-x)^T \Sigma^{-1} (U-x) \right) \frac{1}{2} \exp \left(-\frac{1}{2} (U-x)^T \Sigma^{-1} (U-x) \right) \pi$$

$$\left((U-x)^T \Sigma^{-1} (U-x) \right) \frac{1}{2} \exp \left(-\frac{1}{2} (U-x)^T \Sigma^{-1} (U-x) \right) \pi$$

$$\left((U-x)^T \Sigma^{-1} (U-x) \right) \frac{1}{2} \exp \left(-\frac{1}{2} (U-x)^T \Sigma^{-1} (U-x) \right) \pi$$

$$\left((U^T \Sigma^{-1} U + x^T \Sigma^{-1} x - U^T \Sigma^{-1} x - x^T \Sigma^{-1} U) \right) \frac{1}{2} \exp \left(-\frac{1}{2} (U-x)^T \Sigma^{-1} (U-x) \right) \pi$$

homework_2

February 20, 2025

```
[106]: import pandas as pd
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis as LDA
from sklearn.naive_bayes import MultinomialNB as MNB
from sklearn.metrics import accuracy_score
import math
import numpy as np
from numpy.linalg import qr, norm
import matplotlib.pyplot as plt
from sklearn.mixture import GaussianMixture
import matplotlib.cm as cm
import urllib

[3]: X_train_input = pd.read_csv("20news-bydate/matlab/train.data", delimiter="␣",
    ↪names = ["docIdx", "wordIdx", "freq"],)
y_train_input = pd.read_csv("20news-bydate/matlab/train.label",␣
    ↪names=["labels"])
y_train_input['docIdx'] = y_train_input.index + 1
X_test_input = pd.read_csv("20news-bydate/matlab/test.data", delimiter=" ",␣
    ↪names = ["docIdx", "wordIdx", "freq"],)
y_test_input = pd.read_csv("20news-bydate/matlab/test.label", names=["labels"])
y_test_input['docIdx'] = y_test_input.index + 1

[4]: word_cnt = X_train_input[["wordIdx", "freq"]].groupby(["wordIdx"],␣
    ↪as_index=False).sum().sort_values(by='freq', ascending=False)
word_cnt_filtered = word_cnt[word_cnt['freq']>1000].reset_index()

def preprocessing(X_df, y_df):
    X_df_filtered = X_df.loc[X_df["wordIdx"].isin(word_cnt_filtered.wordIdx)].
    ↪reset_index(drop=True)

    # combine X_* and y_*
    combined_data = X_df_filtered.merge(y_df, on="docIdx", how="inner")
    X = combined_data
    y = combined_data['labels']
    return X,y
```

```
[5]: X_train, y_train = preprocessing(X_train_input, y_train_input)
X_test, y_test = preprocessing(X_test_input, y_test_input)
```

```
[7]: def tfidf_matrix(data, words):
    total_doc = len(data)#.groupby(["wordIdx"]))
    idf = []
    for index, group in data.groupby(["wordIdx"], as_index = False).agg(list).
iterrows():
        idf.append( math.log ( total_doc / len(group.docIdx) ) )
    X = []
    y = []
    tfidf = []
    for index, group in data.groupby(["docIdx"], as_index = False).agg(list).
iterrows():
        x = [0 for i in range(len(words))]
        y.append(group.labels[0])
        for ind in range(len(words)):
            # print(words[ind])
            if words.wordIdx[ind] in group.wordIdx:
                x[ind] = group.freq[group.wordIdx.index(words.wordIdx[ind])]
            # if word[1].wordIdx in group.wordIdx:
            #     x[words.index(word)] = group.freq[group.wordIdx.index(word[1].
iterrows())] / sum(group.freq)
        X.append(x)

    for i in range(len(X)):
        x = []
        total_words_doc = sum(X[i])
        for j in range(len(X[0])):
            x.append(X[i][j] / total_words_doc * idf[j])
        tfidf.append(x)

    # print(y[0])
    return np.array(tfidf), np.array(y)
```

```
[8]: X_train_tfidf, y_train_tfidf = tfidf_matrix(X_train, word_cnt_filtered)
X_test_tfidf, y_test_tfidf = tfidf_matrix(X_test, word_cnt_filtered)
```

0.0.1 1

```
[119]: mnb = MNB()
mnb.fit(X_train_tfidf, y_train_tfidf)
y_pred = mnb.predict(X_test_tfidf)
accuracy_mnb = accuracy_score(y_pred, y_test_tfidf)
print("Accuracy on test data", accuracy_mnb*100)
```

Accuracy on test data 40.60250599840043

0.0.2 2

```
[10]: lda = LDA()  
lda.fit(X_train_tfidf, y_train_tfidf)  
y_pred = lda.predict(X_test_tfidf)
```

```
[11]: accuracy = accuracy_score(y_pred, y_test_tfidf)
```

```
[120]: print("Accuracy on test data", lda.score(X_test_tfidf, y_test_tfidf) * 100)
```

Accuracy on test data 39.53612370034658

0.0.3 3

```
[46]: phi = X_train_tfidf  
phi.shape
```

```
[46]: (11260, 292)
```

```
[47]: def orthogonal_iteration(phi, k, terminate=1e-5):  
    m,n = phi.shape  
    Q = np.random.randn(n, k)  
    Q, R = qr(Q)  
  
    while True:  
        temp = phi @ Q  
        Q_new, R = qr(phi.T @ temp)  
  
        if norm(Q_new - Q) < terminate:  
            break  
  
        Q = Q_new  
  
    return Q
```

```
[48]: theta = orthogonal_iteration(phi, 2)  
theta.shape
```

```
[48]: (292, 2)
```

```
[54]: Y = phi @ theta
```

```
[55]: # plot Y  
Y.shape
```

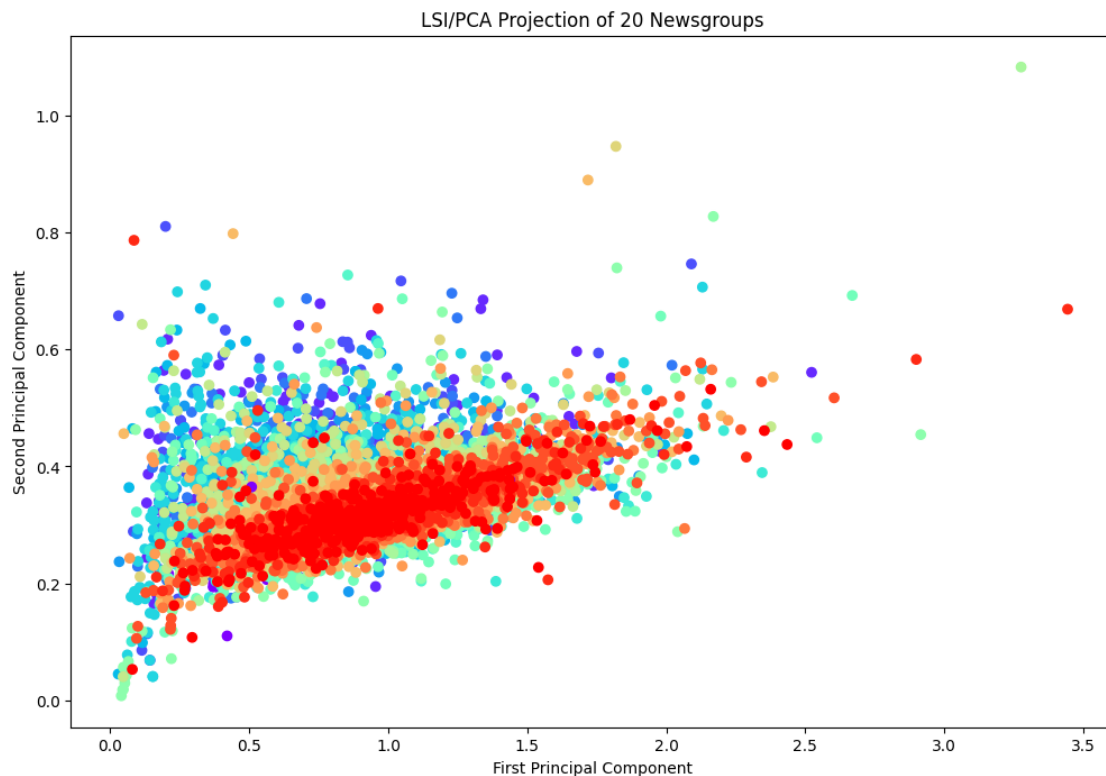
```
[55]: (11260, 2)
```

```
[90]: colors = cm.rainbow(np.linspace(0, 1, 20))  
plt.figure(figsize=(12, 8))
```

```

# for i in range(len(Y)):
#     plt.scatter(Y[i][0], Y[i][1], c=colors[i for j in X_train.
#         ↳groupby(['labels']).agg(set) if i in j])
arr = []
doc_groups = X_train.groupby(['labels']).agg(set)
for i in range(len(Y)):
    flag = 1
    for j, group in doc_groups.iterrows():
        if i+1 in group.docIdx:
            arr.append(j-1)
            flag = 0
            break
    if flag:
        arr.append(0)
plt.scatter(Y[:, 0], Y[:, 1], c=colors[arr])
plt.xlabel('First Principal Component')
plt.ylabel('Second Principal Component')
plt.title('LSI/PCA Projection of 20 Newsgroups')
plt.show()

```



0.0.4 4

```
[94]: phi = orthogonal_iteration(X_train_tfidf, 100)
```

```
[95]: gmm = GaussianMixture(n_components=20, covariance_type='full')
```

```
[96]: gmm.fit(phi)
```

```
[96]: GaussianMixture(n_components=20)
```

```
[97]: means = gmm.means_
```

```
[ ]:
```

```
[98]: Y = phi @ means.T
```

```
[99]: means.shape
```

```
[99]: (20, 100)
```

```
[100]: phi.shape
```

```
[100]: (292, 100)
```

```
[101]: Y.shape
```

```
[101]: (292, 20)
```

```
[115]: url = "http://qwone.com/~jason/20Newsgroups/vocabulary.txt"
file = urllib.request.urlopen(url)

word_index_map = {}
i = 0
for line in file:
    word_index_map[i] = line.decode("utf-8").strip()
    i = i + 1
```

```
[117]: for i in range(20):
        temp = Y[:, i]
        ind1 = temp.argsort()[-10:][::-1]
        print("Top 10 words in cluster", i+1, ":")
        for i in ind1:
            print(word_index_map[i], end=" ")
        print()
```

Top 10 words in cluster 1 :

laurel organizations atrocities san feet francisco box can netcom foundation

Top 10 words in cluster 2 :

and blueprints hannover area secular south bay berlin this internationaler

Top 10 words in cluster 3 :
 name immoralities letters alternate book critiques claus describe well hrsg
 Top 10 words in cluster 4 :
 write newsletter horrors jr islington santa dead atrocities contradictions that
 Top 10 words in cluster 5 :
 archive austin books ny contradictions foote ink critiques miller blueprints
 Top 10 words in cluster 6 :
 for americans wrote promoting claus leaving und post der that
 Top 10 words in cluster 7 :
 december similarity well east per books area atrocities nl this
 Top 10 words in cluster 8 :
 last feet this jr american book well moulded so based
 Top 10 words in cluster 9 :
 hollywood book pp lines ethical james handbook moulded davy conduit
 Top 10 words in cluster 10 :
 religion horrors british thought they short living word be edgar
 Top 10 words in cluster 11 :
 which copying letters com wrote exists informationen many society immoralities
 Top 10 words in cluster 12 :
 plastic this area king containing immoralities pp square monks characters
 Top 10 words in cluster 13 :
 bumper who aah monks claus that horrors short north fax
 Top 10 words in cluster 14 :
 in provoking freethinker one vertrieb absurdities prometheus this books monks
 Top 10 words in cluster 15 :
 addresses humanism berlin kingdom hannover horrors pp word new figmo
 Top 10 words in cluster 16 :
 paraphernalia magazine that exists canyon holloway press lists king ethical
 Top 10 words in cluster 17 :
 dick rationalist that netcom magazine informationen claus islington set national
 Top 10 words in cluster 18 :
 atheism square is zeit miz so norm netcom informationen provoking
 Top 10 words in cluster 19 :
 alt cameron bible passage lion is wrote publish pp one
 Top 10 words in cluster 20 :
 it square provoking various thomas one be pp any ultimate

```
[121]: ans = """
Most of the words in the top 10 are common words like can, and, it, is
but some words like hollywood in cluster 10 or organisations in cluster 1
show that the cluster might have new articles about hollywood or organisations
respectively.
"""
print(ans)
```

Most of the words in the top 10 are common words like can, and, it, is
 but some words like hollywood in cluster 10 or organisations in cluster 1

show that the cluster might have new articles about hollywood or organisations respectively.

[]: